

Deploying EDA Workloads On Public Cloud

Melvin Cardozo



Topics

- Cloud Motivations
- EDA compute workloads on Traditional, Cloud-First, Hybrid infrastructure
- Cloud Native Infrastructure and Terminology
- Customer Cloud Challenges
- Current Cloud Status For Synopsys Tools
- Case Study 1 - VCS
- Case Study 2 – PrimeTime, ICD2

Cloud Motivations

Cloud Participants in EDA

Motivation:


- Optimize the use of HW farm spend
- Use the best tools, methodology

New Motivation:

- Short term project-specific needs
- Improve TTR, QoR
- Managed Services
- New business (ARM, AMD, nvda, xlnx)

New Motivation:

- Make money from HW & managed services
- EDA compute is a big target (with specific infra requirements)



Motivation:


- Best / fastest engines
- Take advantage of h/w

New Motivation:

- Use elastic & unlimited scale to improve product vs competition
- Make money off software in this segment
- Make money off special purpose hardware in this segment

Astera Labs: VCS on Amazon Cloud

Actual Usage with Hourly License Model
87,078 License-hours over 28 days



28 days with unlimited VCS licenses

Equivalent Usage with TSL Model
87,078 License-hours over 227 days



227 days with 16 VCS licenses

- Customer had only 16 VCS licenses available under TSL model
- Monitored real project use with “unlimited” cloud access to VCS – 87K license-hours used in 1 month
- Cloud access effectively delivered 8x faster TAT and 6+ months shift-left

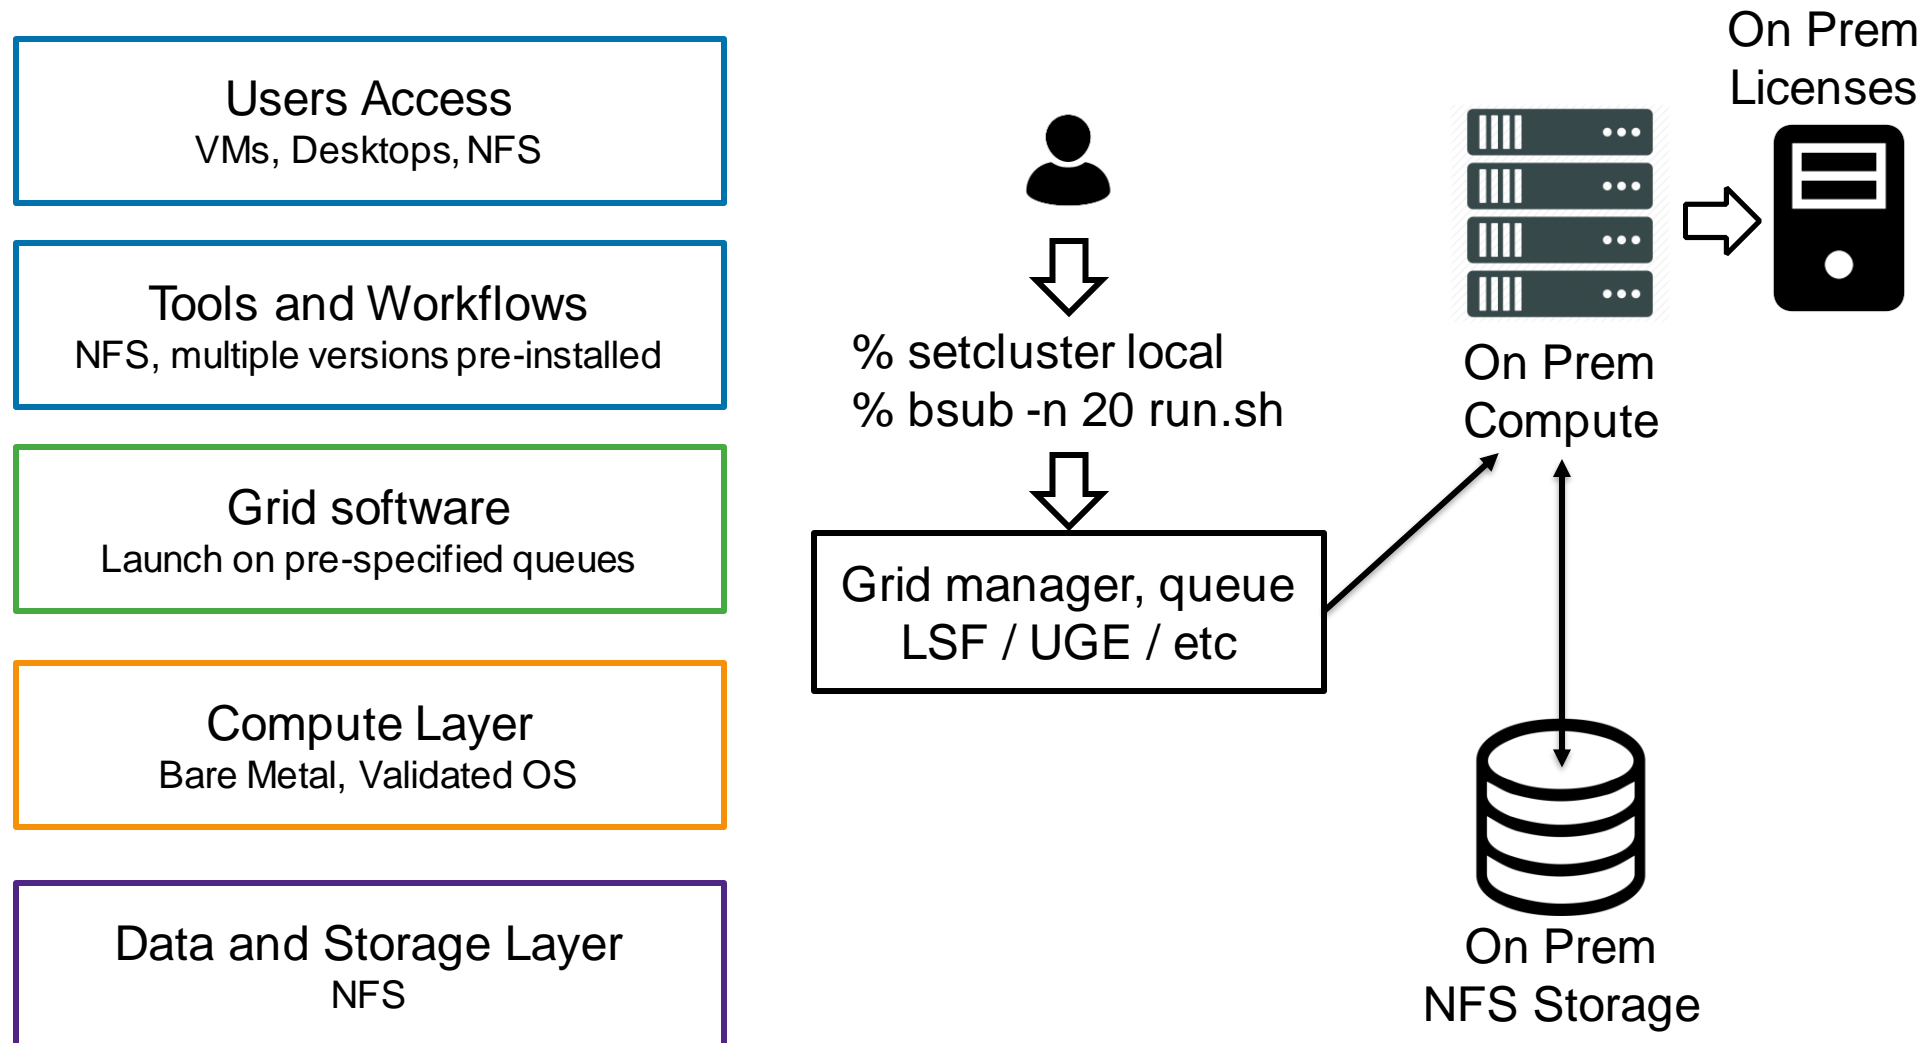
Cloud can
enable huge
shift-left!

Traditional, Cloud-First & Hybrid Cloud

Verification on Cloud Segments

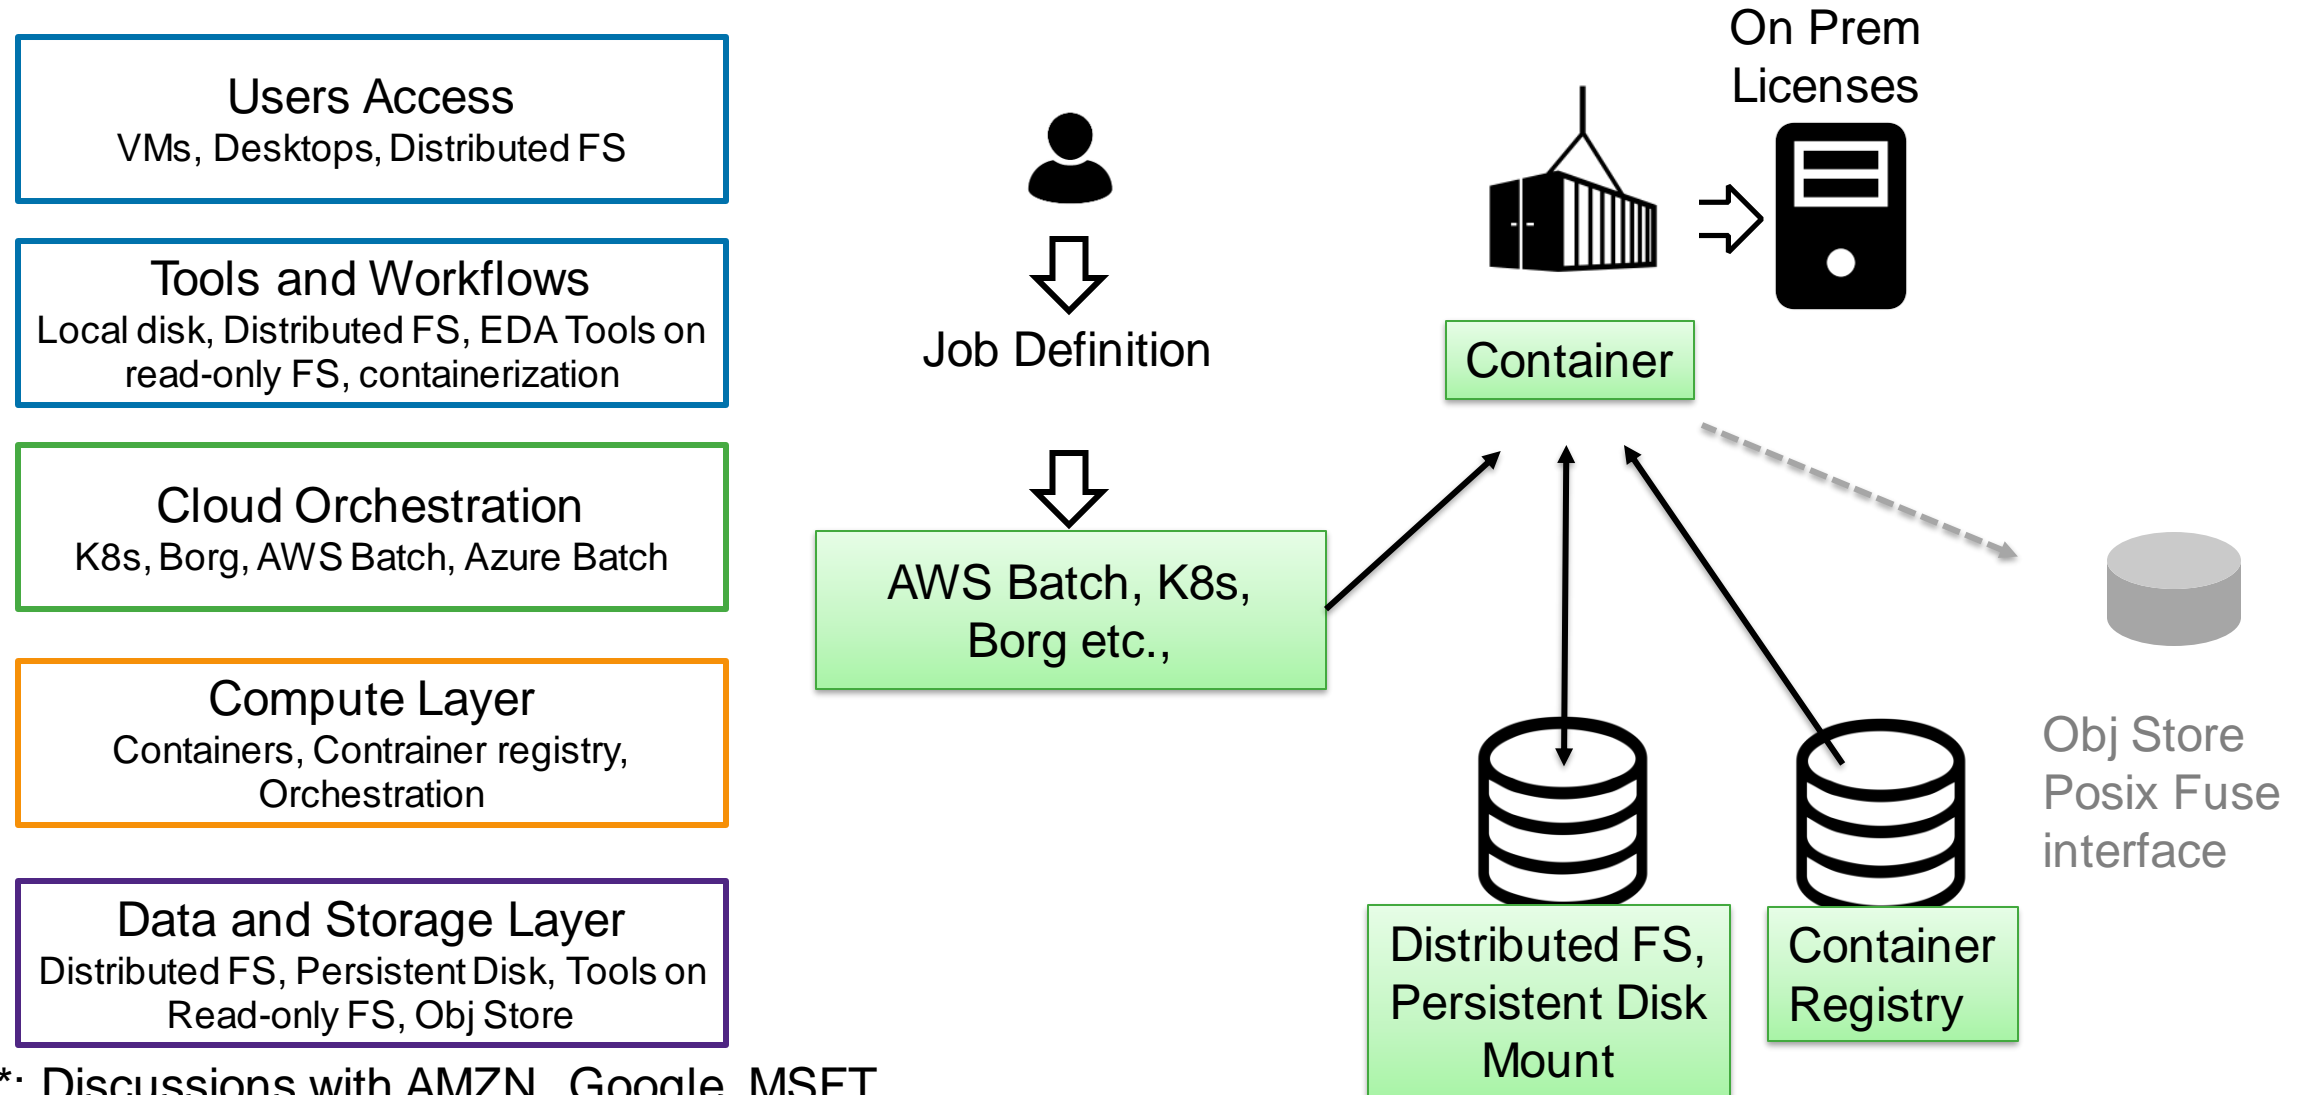
	Large Semi	Mega Disruptors	Startup Semi
Description	Traditional chip design companies with on-premise resources for EDA workloads	Companies which provide web services, products, platforms & infrastructure; moving into value added chip designs	Small, new fabless semiconductor companies looking to bring initial products to market quickly
Large, existing IT team and compute infra	Yes	Yes	No
Large, legacy EDA infrastructure	Yes	No	No
Existing compute infra designed for cloud-native workloads	No	Yes	No
Require “burst” access for peak or project use	Yes	No	No
CAPEX-constrained	No	No	Yes
Target cloud usage model	Hybrid for peak	Private	Public
EDA & related IT infrastructure spend	\$\$\$	\$\$	\$
Examples	AMD, Intel, Nvidia, Qualcomm, Samsung, Broadcom, ARM, Xilinx	Google, Amazon, Microsoft, Facebook, Baidu, Alibaba	Astera Labs, Groc, Cerebras, Esperanto, Sambanova, Graphcore

EDA workloads: Traditional Architecture



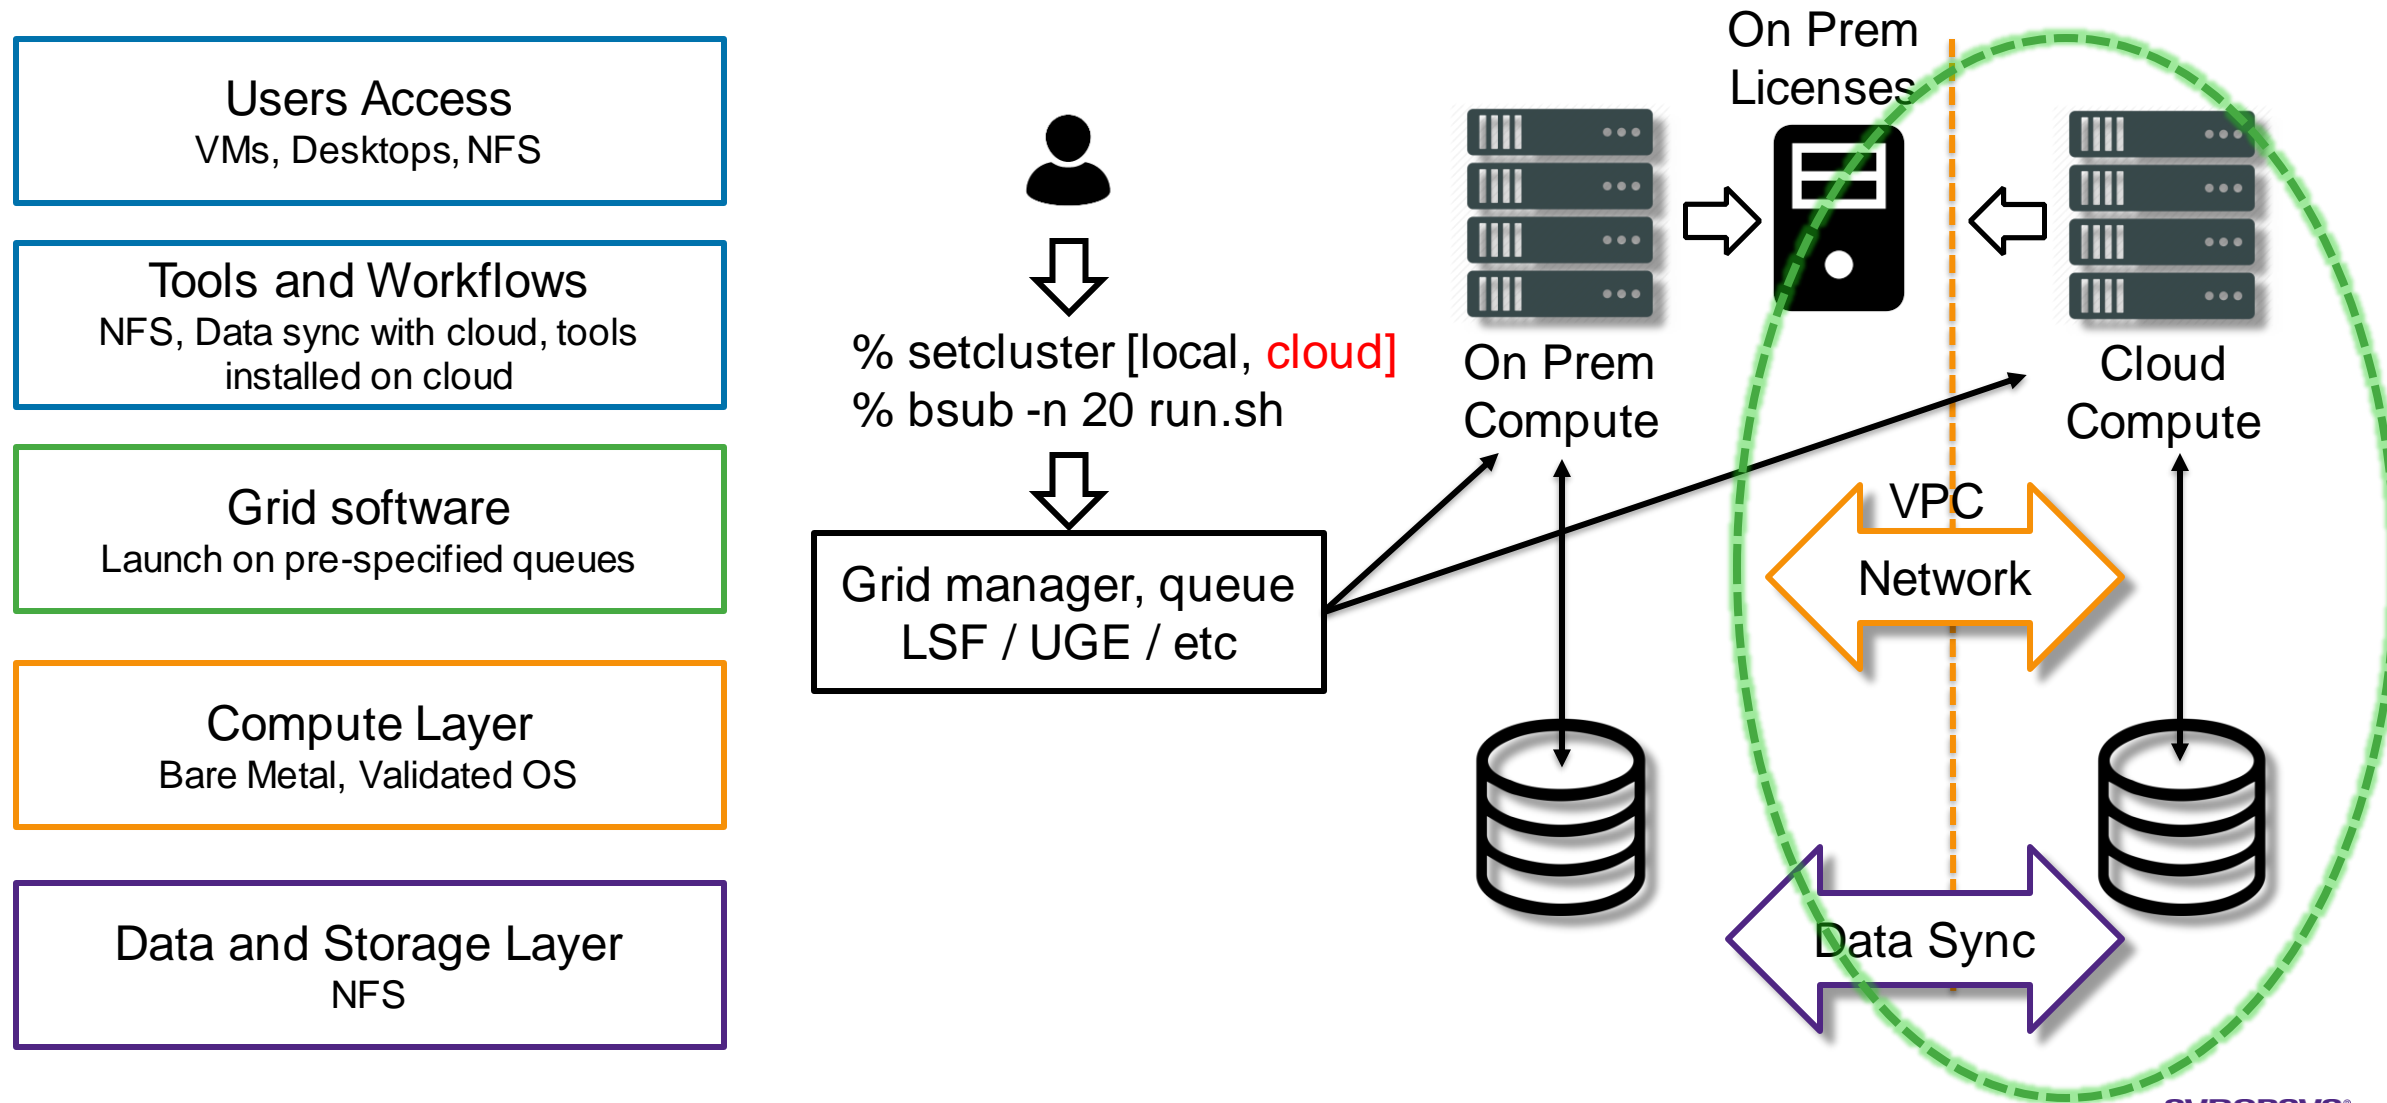
EDA workloads: Cloud First Architecture

One of several options, not including security, authentication, authorization



*: Discussions with AMZN, Google, MSFT

Hybrid cloud for traditional EDA users



Cloud Native Infrastructure And Terminology

Amazon Cloud Offerings



Analytics



Application Integration



AR & VR



AWS Cost Management



Blockchain



Business Applications



Compute



Customer Engagement



Database



Developer Tools



End User Computing



Game Tech



Internet of Things



Machine Learning



Management & Governance



Media Services



Migration & Transfer



Mobile



Networking & Content
Delivery



Robotics



Satellite



Security, Identity &
Compliance

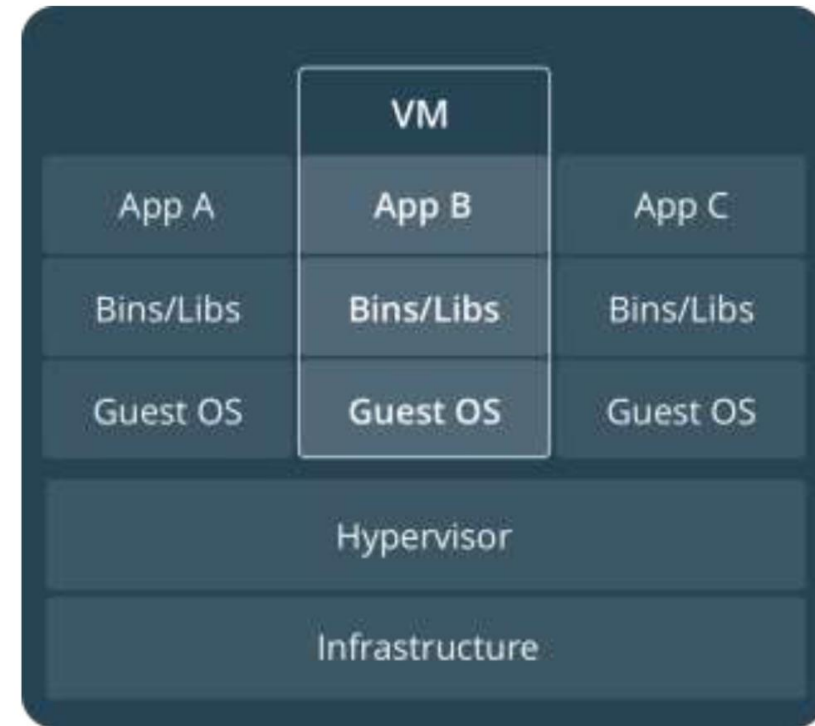


Storage

Containers v/s VMs

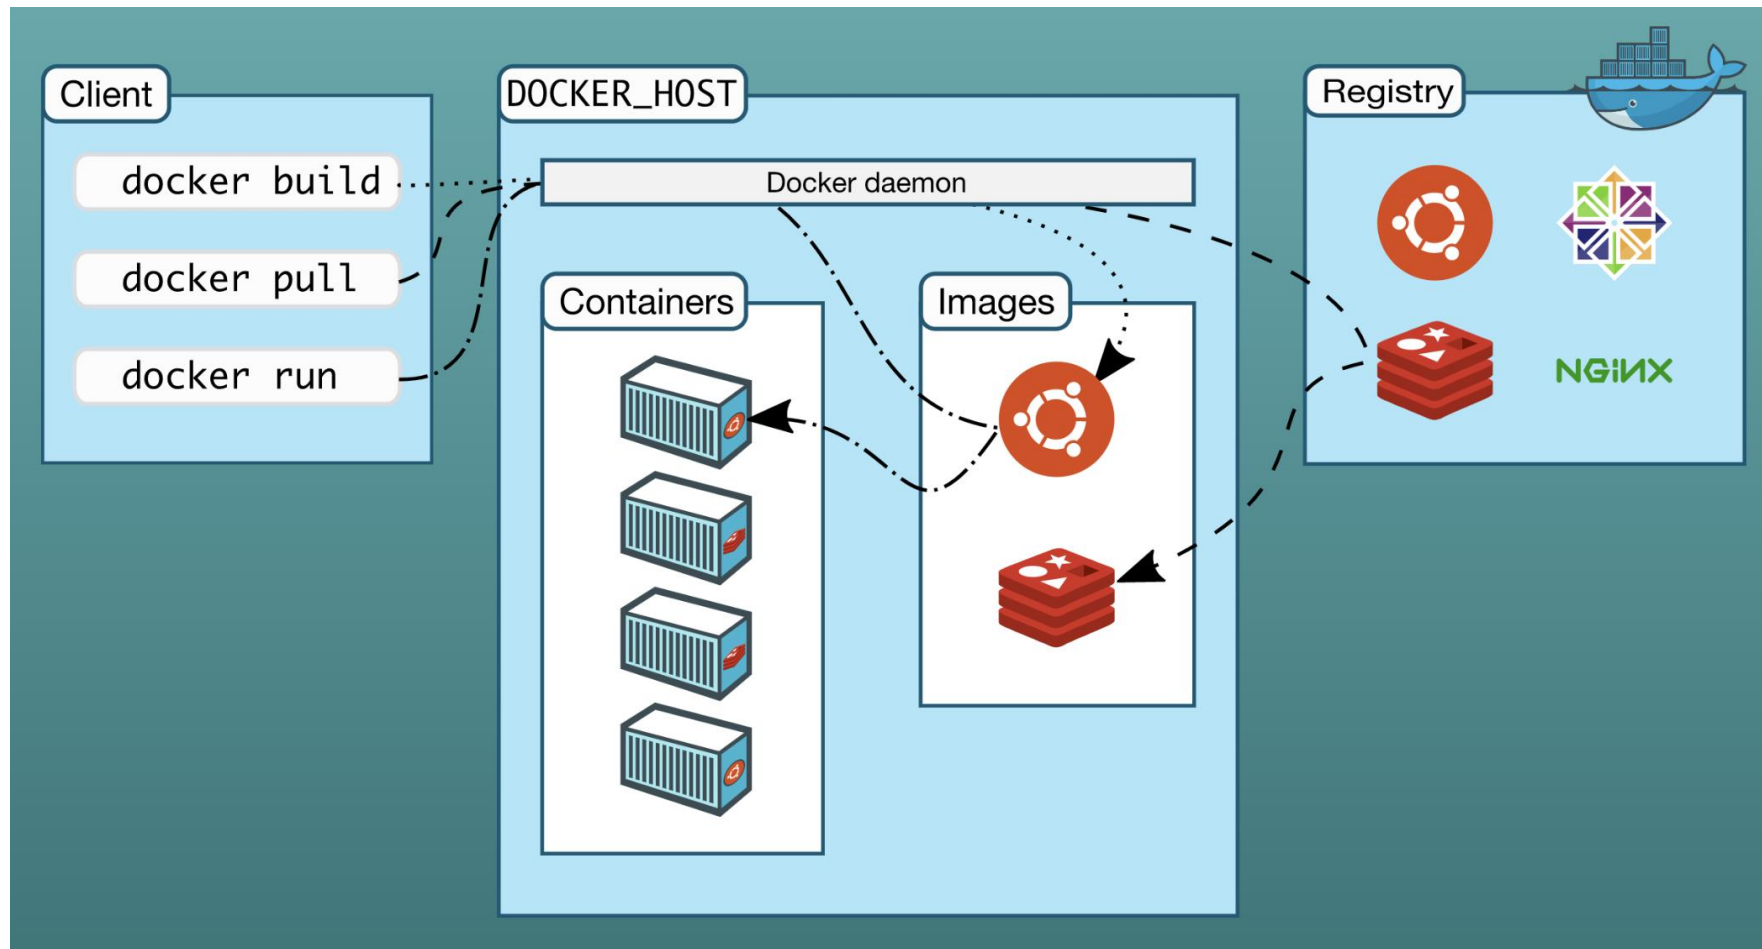


Containers are an app
level construct



VMs are an infrastructure level
construct to turn one machine
into many servers

Docker Architecture



Docker Image v/s Container

- Docker Image
 - Docker images are **read-only** templates from which Docker containers are launched.
 - Each image consists of a series of **layers**.
 - Docker makes use of **Union file systems** to combine these **layers** into a single image.
- Docker Container
 - **Container** is built from an **image**.
 - A container consists of an operating system, user-added **files**, and **meta-data**.
 - That image tells Docker what the container **holds**, what process to **run** when the container is launched, and a variety of other **configuration** data.
 - The Docker image is **read-only**.
 - When Docker runs a container from an image, it adds a read-write layer on top of the image (using a union file system as we saw earlier) in which your application can then run.

Docker Image Registry

- Images are stored locally, but can be pulled from an Image registry
- The Registry is a server that stores and lets you distribute Docker images.

Dockerfile

```
FROM centos:7

ARG INSTALL_DIR

# Install the required dependencies
RUN set -ex \
    && yum makecache fast \
    && yum -y update \
    && yum -y install epel-release \
    && yum -y install mariadb-server \
    && yum clean all \
    && rm -rf /var/cache/yum

# Create /usr/local/bin
RUN set -ex \
    && mkdir -p /usr/local/bin

# Install slurm
COPY $INSTALL_DIR/bin $INSTALL_DIR/bin
COPY $INSTALL_DIR/include $INSTALL_DIR/include
COPY $INSTALL_DIR/lib64 $INSTALL_DIR/lib64
COPY $INSTALL_DIR/sbin $INSTALL_DIR/sbin
COPY $INSTALL_DIR/share $INSTALL_DIR/share

# Create /etc/slurm
RUN set -ex \
    && mkdir -p /etc/slurm

# Install the config files
COPY $CONFIG_BASE/etc/slurm/slurmdbd.conf $INSTALL_DIR/etc/slurm/slurmdbd.conf
COPY $CONFIG_BASE/deploy/DATABASE/docker-entrypoint.sh /usr/local/bin/docker-entrypoint.sh
RUN sed -i -r "s#<INSTALL_DIR>#$INSTALL_DIR#g" /usr/local/bin/docker-entrypoint.sh
ENTRYPOINT ["/usr/local/bin/docker-entrypoint.sh"]

CMD ["slurmdbd"]
```

Docker Compose

```
version: '3'

services:
  mysql:
    image: mysql
    hostname: mysql
    container_name: mysql
    environment:
      MYSQL_DATABASE: slurm_acct_db
      MYSQL_ROOT_PASSWORD: root
      MYSQL_USER: slurm
      MYSQL_PASSWORD: password
    volumes:
      - ./database:/var/lib/mysql
    ports:
      - 3306:3306
```

```
slurmdbd:
  image: slurmdbd:19.05
  command: [slurmdbd]
  container_name: slurmdbd
  hostname: slurmdbd
  volumes:
    - var_log_slurm:/var/log/slurm
  ports:
    - 6819:6819
  depends_on:
    - mysql

volumes:
  var_log_slurm:
```

Docker Build & Run

```
docker build -t slurmdbd:19.05 \  
  --build-arg CONFIG_BASE=$CONFIG_BASE \  
  --build-arg INSTALL_DIR=$INSTALL_DIR \  
  --build-arg TARGET_INSTALL_DIR=$TARGET_INSTALL_DIR \  
  --build-arg USER=$USER \  
  --build-arg GROUP=$GROUP \  
  --build-arg UID=$UID \  
  --build-arg GID=$GID \  
  -f Dockerfile \  
  ../../..
```

```
docker volume create var_log_slurm  
docker volume create var_lib_mysql
```

```
docker run -d \  
  --name mysql \  
  --hostname mysql \  
  --env MYSQL_DATABASE=slurm_acct_db \  
  --env MYSQL_ROOT_PASSWORD=root \  
  --env MYSQL_USER=slurm \  
  --env MYSQL_PASSWORD=password \  
  --volume var_lib_mysql:/var/lib/mysql \  
  --publish 3306:3306 \  
  mysql:5.7
```

```
docker run -d \  
  --name slurmdbd \  
  --hostname slurmdbd \  
  --volume var_log_slurm:/var/log/slurm \  
  --publish 6819:6819 \  
  slurmdbd:19.05
```

Object Store

Object Storage Systems Characteristics

- Data is stored as individual objects with a unique identifier
- Flat addressing scheme that allows for greater scalability
- Multi-tenant
- Usually software-based that runs on commodity hardware
- Capable of scaling to 100s of Petabytes
- Don't use RAID but instead Replication and/or Erasure Coding
 - At PBs scale RAID has very long rebuild times
- Access over RESTful API over HTTP, which is a great fit for cloud and mobile applications
 - Amazon S3, Swift and CDMI API

Object Storage is a good fit for

- Unstructured data workloads
- Capacity requirements beyond 100s of Terabytes
- Distributed access to content
- Data archiving: documents, email, backups
- Storage for photos, videos, virtual machines images
- Need for granular security and multi-tenancy
- Need for automation, management, monitoring and reporting tools
- Non-high performance applications



Consumer Activity
(Events, GPS, WIFI)



Device Tracking and Logs
(Event, Configuration, Usage, Performance,)



Social media

Object Store APIs

Container/Bucket

GET	/v1/ {account} / {container}	Show container details and list objects
PUT	/v1/ {account} / {container}	Create container
POST	/v1/ {account} / {container}	Create, update, or delete container metadata
HEAD	/v1/ {account} / {container}	Show container metadata
DELETE	/v1/ {account} / {container}	Delete container

Objects

GET	/v1/ {account} / {container} / {object}	Get object content and metadata
PUT	/v1/ {account} / {container} / {object}	Create or replace object
COPY	/v1/ {account} / {container} / {object}	Copy object
DELETE	/v1/ {account} / {container} / {object}	Delete object
HEAD	/v1/ {account} / {container} / {object}	Show object metadata
POST	/v1/ {account} / {container} / {object}	Create or update object metadata

Batch Schedulers

Introducing AWS Batch



Fully Managed

No software to install or servers to manage. AWS Batch provisions, manages, and scales your infrastructure



Integrated with AWS

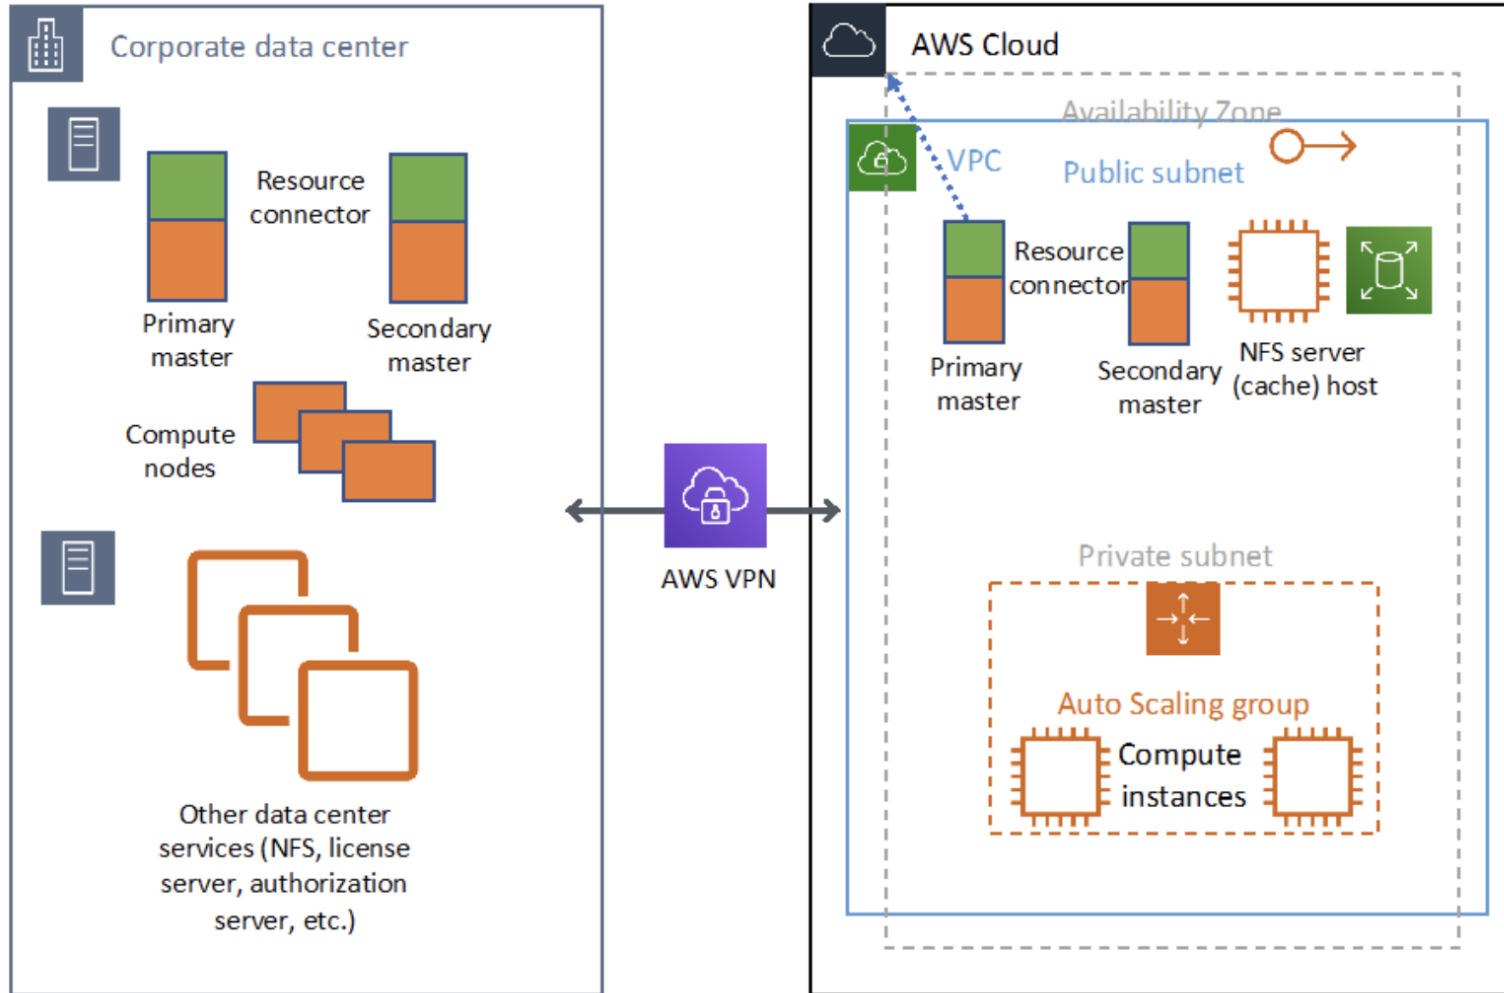
Natively integrated with the AWS Platform, AWS Batch jobs can easily and securely interact with services such as Amazon S3, DynamoDB, and Rekognition



Cost-optimized Resource Provisioning

AWS Batch automatically provisions compute resources tailored to the needs of your jobs using Amazon EC2 and EC2 Spot

LSF Connect



Cloud Pricing - AWS

<50%
Utilization

Instance Types	Cores	Memory (MB)	Disk	Description	OnDemand Price	Spot Price
t3.nano – t3.2xlarge	1 - 4	0.5 - 32	EBS	0.5GB/core	\$0.005 - \$0.332	32%
c5.large – cg.24xlarge	1 - 48	4 - 92	EBS	4GB/core	\$0.085 – \$4.08	38%
c5d.large – c5d.18xlarge	1 - 36	4 - 144	50G – 1.8T	4GB/core, SSD	\$0.096 - \$3.456	33%
m5.large – m5.24xlarge	1 - 48	8 - 345	EBS	8GB/core	\$0.096 - \$4.608	35%
m5d.large – m5d.24xlarge	1 - 48	8 - 345	75G – 3.6T	8GB/core, SSD	\$0.113 - \$5.424	30%
r5.large – r5.24xlarge	1 - 48	16 - 768	EBS	16GB/core	\$0.126 - \$6.048	28%
r5d.large – r5d.24xlarge	1 - 48	16 - 768	75G – 3.6T	16GB/core	\$0.144 - \$6.912	25%
i3.large – i3.16xlarge	1 - 32	16 - 512	475 – 15T	16GB/core, SSD	\$0.156 - \$4.992	30%
f1.2xlarge – f1.16xlarge	4 - 32	122 - 976	470 – 4T	1 – 8 FPGAs	\$1.65 - \$13.20	30%
a1.medium – a1.4xlarge	0.5 - 8	2 - 32	EBS	4GB/core, ARM	\$0.025 - \$0.408	33%
g3.4xlarge - g3.16xlarge	8 - 32	47 - 188	EBS	6GB/core, GPU	\$1.14 - \$4.56	30%
p3.2xlarge – p3.16xlarge	8 - 64	26 - 188	EBS	3GB/core, GPU	\$3.06 - \$24.48	30%

AWS EC2 Pricing Models

Instance purchasing option	Risk	Cost	Features
On-demand	Low	High	<ul style="list-style-type: none">• Pay, by the second, for the instances that you launch.
Reserved	Low	Medium	<ul style="list-style-type: none">• Dedicated compute, paid for up-front
Spot	High	Low	<ul style="list-style-type: none">• Spare compute at steep discounts• Spot Instances can be interrupted by EC2 with two minutes of notification when EC2 needs the capacity back.

AWS WhitePaper

- Optimizing Electronic Design Automation (EDA) Workflows on AWS

<https://d1.awsstatic.com/whitepapers/optimizing-electronic-design-automation-eda-workflows-on-aws.pdf>

Customer Cloud Motivations & Challenges

Cloud Challenges For EDA



Cloud is more expensive than on-premise dedicated compute. Network data egress is expensive. Spot instances on cloud comes close (*but still higher*) than on-premise. Cost management, control and allocation is absolutely necessary.



Scalability in HW enables shift-left. EDA license models are the gating factor in flexibility. Astera (startup) was a proof point for shift left with metered licensing, with win-win for Synopsys.



Data transfer to and from cloud is a bottleneck and cost in hybrid and data-intensive flows. Data management and sync with cloud is a work in progress.



Grid scheduler issues for ephemeral instances and / or workload prediction strategy. Data analysis and debug can become bottleneck.



For **startups, cloud is the preferred** solution. Looking for help and qualified flows. Cloud setup still takes too much time (weeks) to start, with too many options.

Cloud security not mentioned as **a primary issue**. That seems past us.

Scale

Customer Needs

- HPC Scaling
 - Scale a single application run across hundreds' of cores (distributed system)
 - To reduce time
 - Handle large designs
 - Elastic scaling
- HTC Scaling
 - Scaling for peak usage to reduce overall turnaround time
- Workflow management
- Non-standard hardware
- Emulation hardware Scaling

Product Needs

- Products that can elastically scale their resource requirements with minimal communication overhead
- Handle non-standard hardware for cost/performance benefits
 - AMD
 - ARM
 - Cray (Azure)
 - ENI (Enhanced network interface for low latency communication)
- Host ZeBu/HAPS on the cloud

Compute Cost

Customer Needs

- Cloud vendors are pushing EDA towards spot/preemptible instances to bring costs in line with on-premise costs
- Spot/Preemptible instances can be terminated with a very short notice: 30-120s
- Spot/Preemptible instances availability and pricing is dynamic and different for different instance types
- Management interface for putting cost restrictions

Product Needs

- Products can support checkpoint/restore
 - One short
 - Periodic saves
- Products can change requirements dynamically based on available compute resources
 - CPU
 - Memory
- Products can be compiled against a specific target

Hybrid File Access

Customer Needs

- Jobs could run either on-premise or on the cloud depending on hardware availability
 - Typically not known till the application is ready to run
- Keeping data in-sync between on-premise and cloud is a challenge
- Cost of transferring data from the cloud to on-premise is a big factor in determining what can be run on the cloud

Cloud Vendors Provide

- Cloud hosted NetApp
 - Seamless sync between on-prem and cloud
 - No transfer cost (both directions)

Product Needs

- Products can generate executables that can run both on-premise and the cloud
- Packing all the job's dependencies into a container significantly reduces the data sync problem
- Ability to create summarized data and merge sets of summarized data
- Debug/Analysis tools that can be run remotely (close to the data)
- Co-locate emulation hardware on the public cloud

Startups

Customer Needs

- Infrastructure setup is a big issue
 - No dedicated IT expertise
- EDA tool setup is an overhead

Cloud Vendors Provide

- Documents on how to setup an environment
 - Too many options for each requirement
- Visualization and Management
 - e.g Nice DCV

Product Needs

- Startup kit for bringing up a verification optimized infrastructure on each cloud provider
- Prebuilt environments (or images) on each cloud provider with our EDA tool bundles pre-installed
 - Execution Manager
 - ML Platform
 - License Server
 - Different Tool Versions

Current Cloud Status For Synopsys Tools

Becoming Cloud Native

Ready

Feature	
OS Compatibility	Does the tool work without issues on cloud vendor OS (on a supported QSC machine)?
License/Metering access	Can we host the license keys / checkout on Cloud ?
Scale	How much compute capacity can the tool use ?
IO dependency	How much is the tool dependent on NFS ?

Optimized

Distributed computing	Does the tool support usage of large number of cloud nodes ?
Cloud data storage models	Does the tool support cloud based data storage (like Key-Value storage, not just NFS)?
Cloud job scheduling	Does the tool run under schedulers favored by cloud vendors - like AWS Batch, PBS ?
Flow/Data optimization	Is data transfer optimized ? Do tools work in coordination on cloud? Do the runs support burst/hybrid modes of running?

Native

Cloud native packaging	Does the tool use cloud Container technology to bypass Cloud OS compatibility issues ?
Cloud native data usage	Does the tool support cloud data primitives (like S3, Object Storage) ?
Application checkpointing	Does the tool work well when jobs are halted and resumed at will ?
Cloud HPC support	Can the tool make use of high performance components on cloud (Infiniband, Lustre) ?
Elasticity	Can the tool dynamically scale its demand and throughput using more/less capacity ?
Application Security	Does the tool support cloud native security like data encryption and application hardening ?
License scaling	Does the tool optimize checkouts for cloud scale? Does it allow metering in the order of hours?

VG – Cloud Matrix

● Complete. ⊗ Not Available
N/A Not Applicable
 In progress/ partially tested
● Low ● High Dependency Levels

Ready

Feature	VCS/VIP	Static	Formal	Verdi	Wattson	ZeBu	HAPS	Hawk	FPGA	VP
OS Compatibility	●	●	●	●	●	●	●	●	●	●
License/Metering Access	●	●	●	●	●	●	●	●	●	●
Scale	x,000,000	x,00	x000	x,000,000	□	x,000	x,00	x,00	x,000	●
IO dependency	●	●	●	●	●	●	●	●	●	●







Optimized

































































































Distributed computing	●	□	●	●	●	●	●	●	●	N/A
Cloud data storage models	●	●	⊗	□	●	⊗	⊗	●	⊗	⊗
Cloud job scheduling	●	●	●	⊗	□	●	●	●	⊗	⊗

Native

Cloud native packaging										
• Singularity / DP (in-tool)	●	□	⊗	⊗	□	N/A	N/A	●	⊗	⊗
• Docker / Kubernetes	●	□	⊗	⊗	□	N/A	N/A	●	⊗	□
Cloud native data usage	⊗	⊗	⊗	□	●	⊗	⊗	⊗	⊗	⊗
Application checkpointing	●	●	⊗	⊗	□	□	□	●	□	⊗
Cloud HPC support	●	⊗	⊗	⊗	●	⊗	⊗	●	⊗	N/A
Elasticity	●	⊗	⊗	●	□	●	●	●	●	●

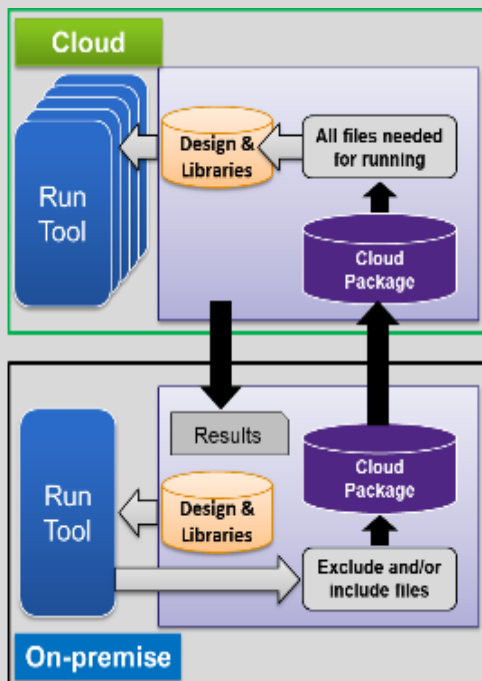
SNPS - Cloud Matrix

 Complete.
  Not Available
 Not Applicable
 In progress / partially tested
 Low  High Dependency Levels


		ICV	Primetime	StarRC	Tetramax	VCS	SiS	Proteus	CATS
Ready	OS Compatibility								
	License Accessibility								
	Scale	x,00	x,00	x,00	x,000	x,000,000	x,00,000	x0,000	x0,000
	NFS Dependency								
Optimized	CDPL Integration								
	CDSL Integration								
	Alternate DRM Support								
Native	Containers								
	• Singularity / DP (in-tool)								
	• Docker / Kubernetes								
	Object Storage								
	Checkpointing								
	HPC Infrastructure								
	Dynamic Scale								

Cloud Building Blocks

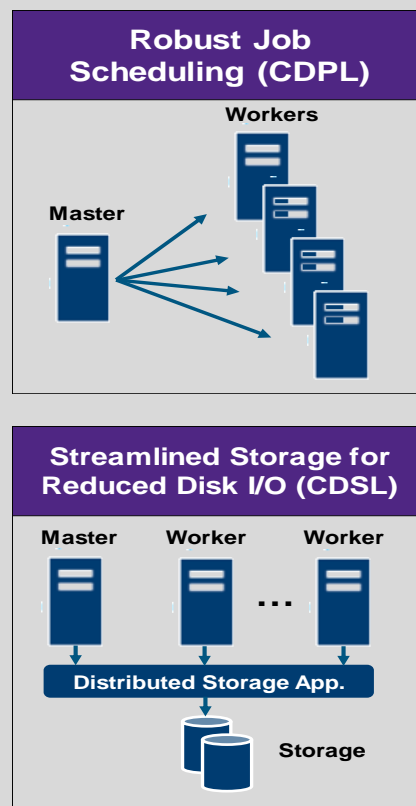
Cloud Transfer Packager



Cloud Tool Containers

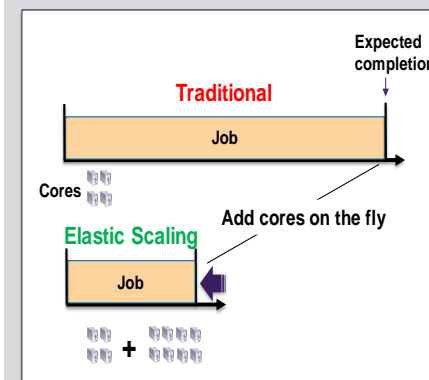


Distributed-Workload Manager

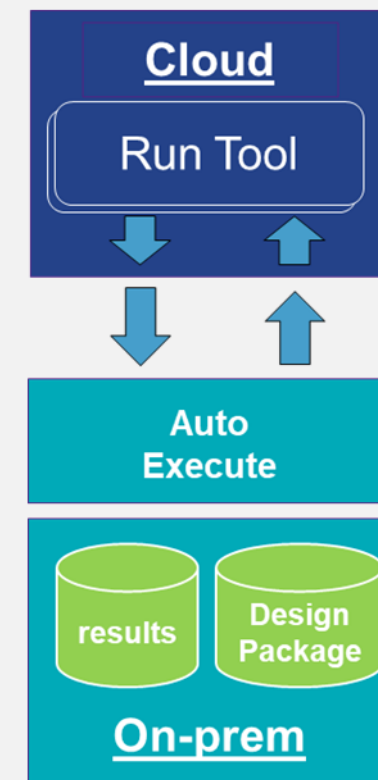


Elastic Scalability

Adds Cores on the fly



Hybrid Connector - \$cloud



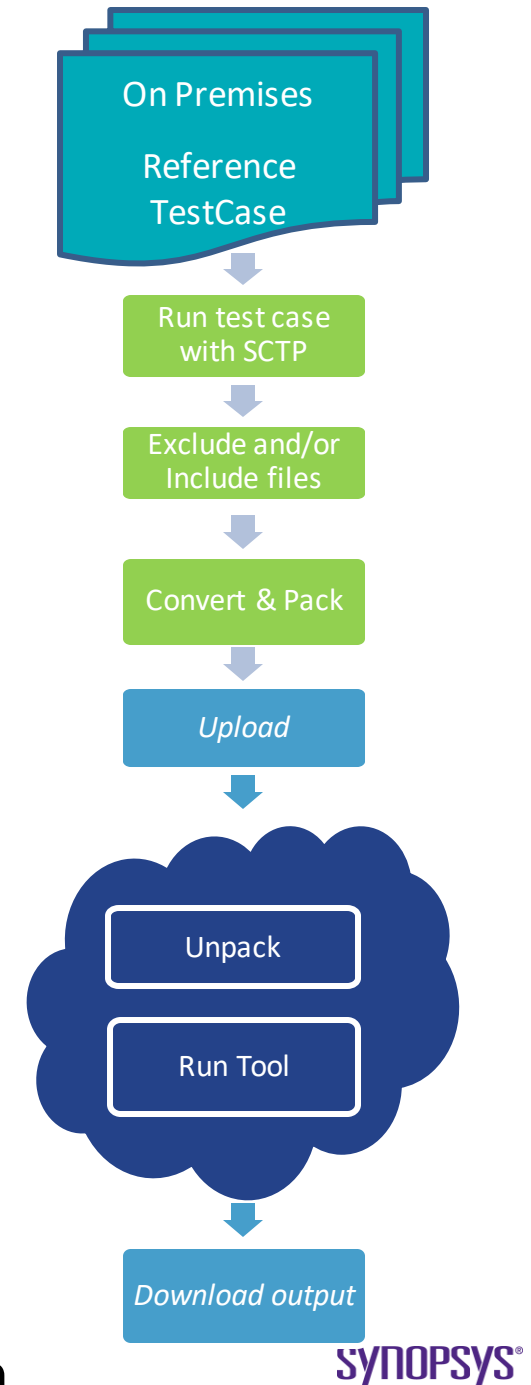
SCTP : Synopsys Cloud Transfer Packager

Product Integration Guide : <https://sweportal/sde/sctp/Pages/Product-Integration.aspx>

DEFINITION : Create & Xfer packages to & from premises to cloud
/depot/tools/sctp/sctp

Key features


- ✓ Standalone / Product independent
- ✓ Generate/Pack & Unpack full data dependency tree*
 - Multithreaded execution
- ✓ Maintains environmental namespace on-prem to cloud
- ✓ MT & Distributed support
- ✓ Reduces upload files sizes
 - Comprehensive exclusion rule set
- ✓ Full Incremental functionality / overlay
 - Upload & download
- ✓ Package creation for Hybrid model execution (-cloud)



CDPL: build complex distributed flows in our tools

Simple,
uncoupled
operations

- Highly parallel operations




Simple,
time
dimension
flows

- Single-lane sequenced fast operations




Domain
specific
tasks

- Operations with dynamic dependencies



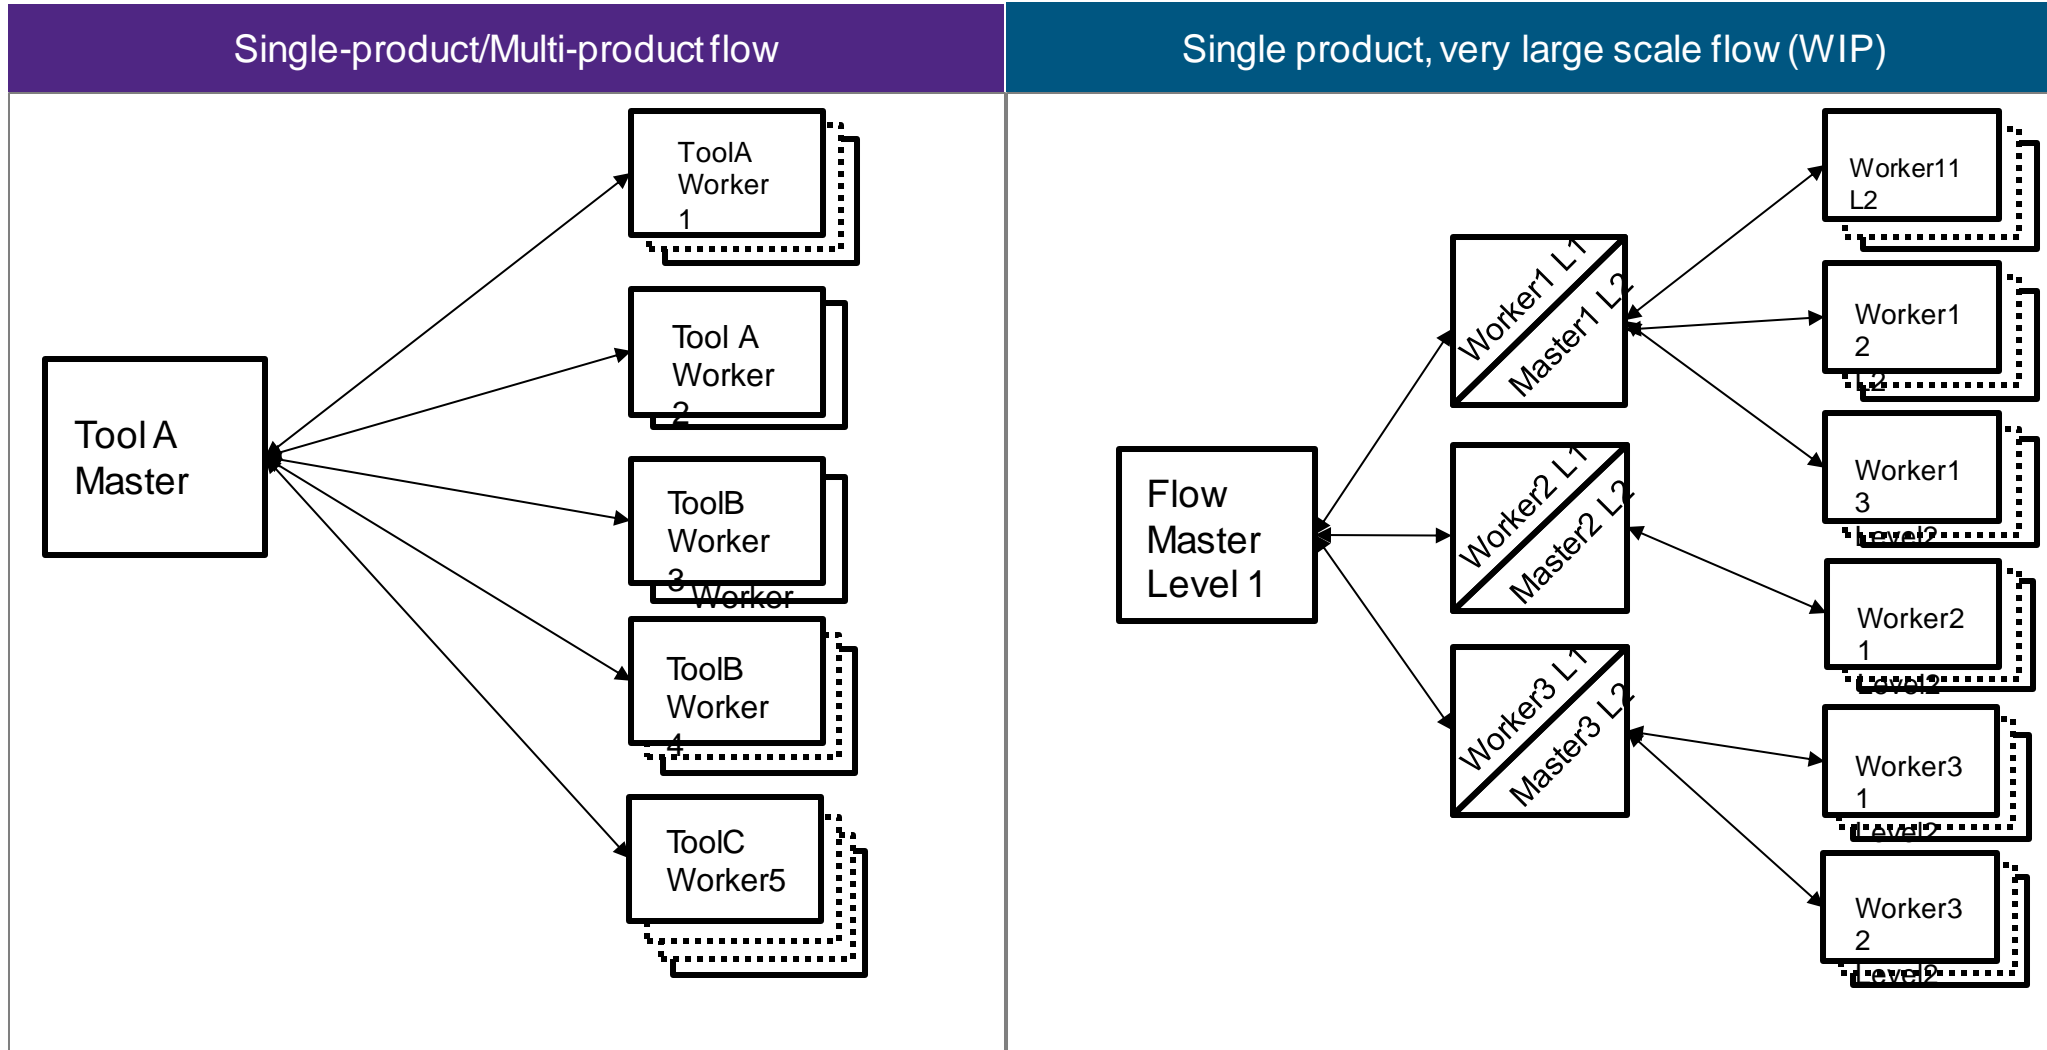
Phases,
Checkpoint
s

- Batched/staged operations

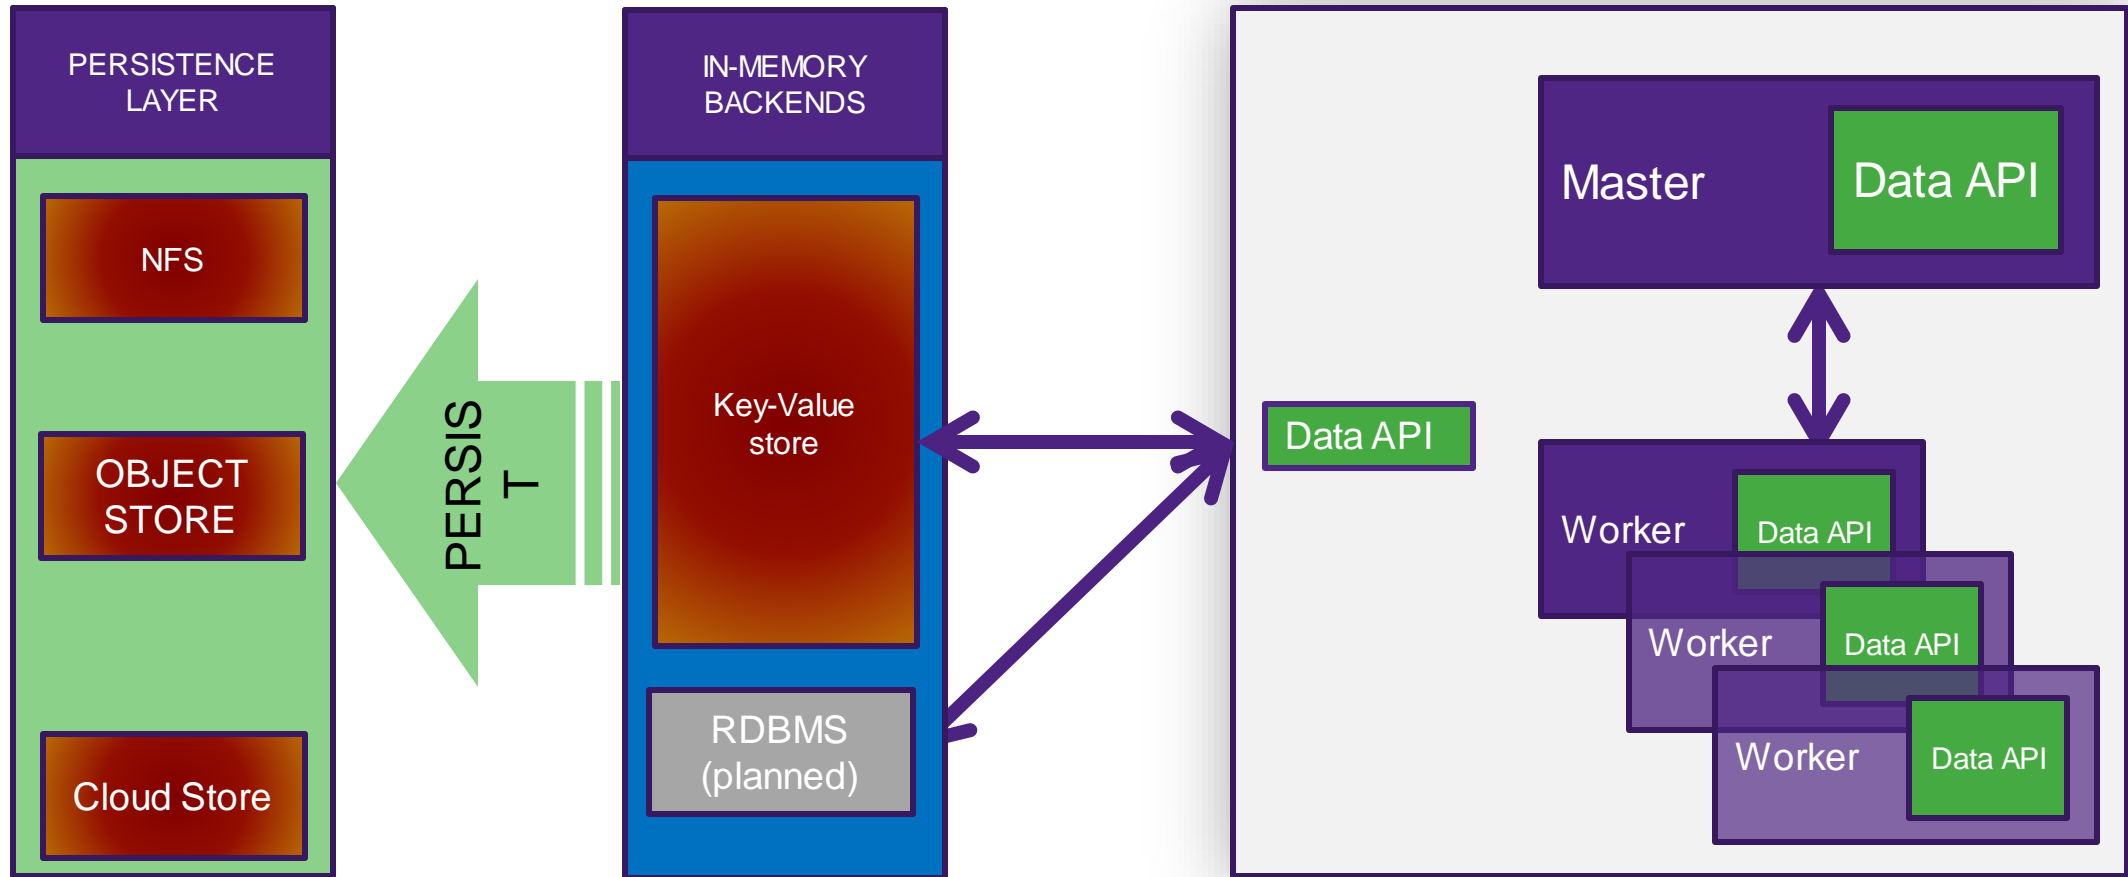


Examples of applications (today, and in the works)

Large and Complex DP possibilities

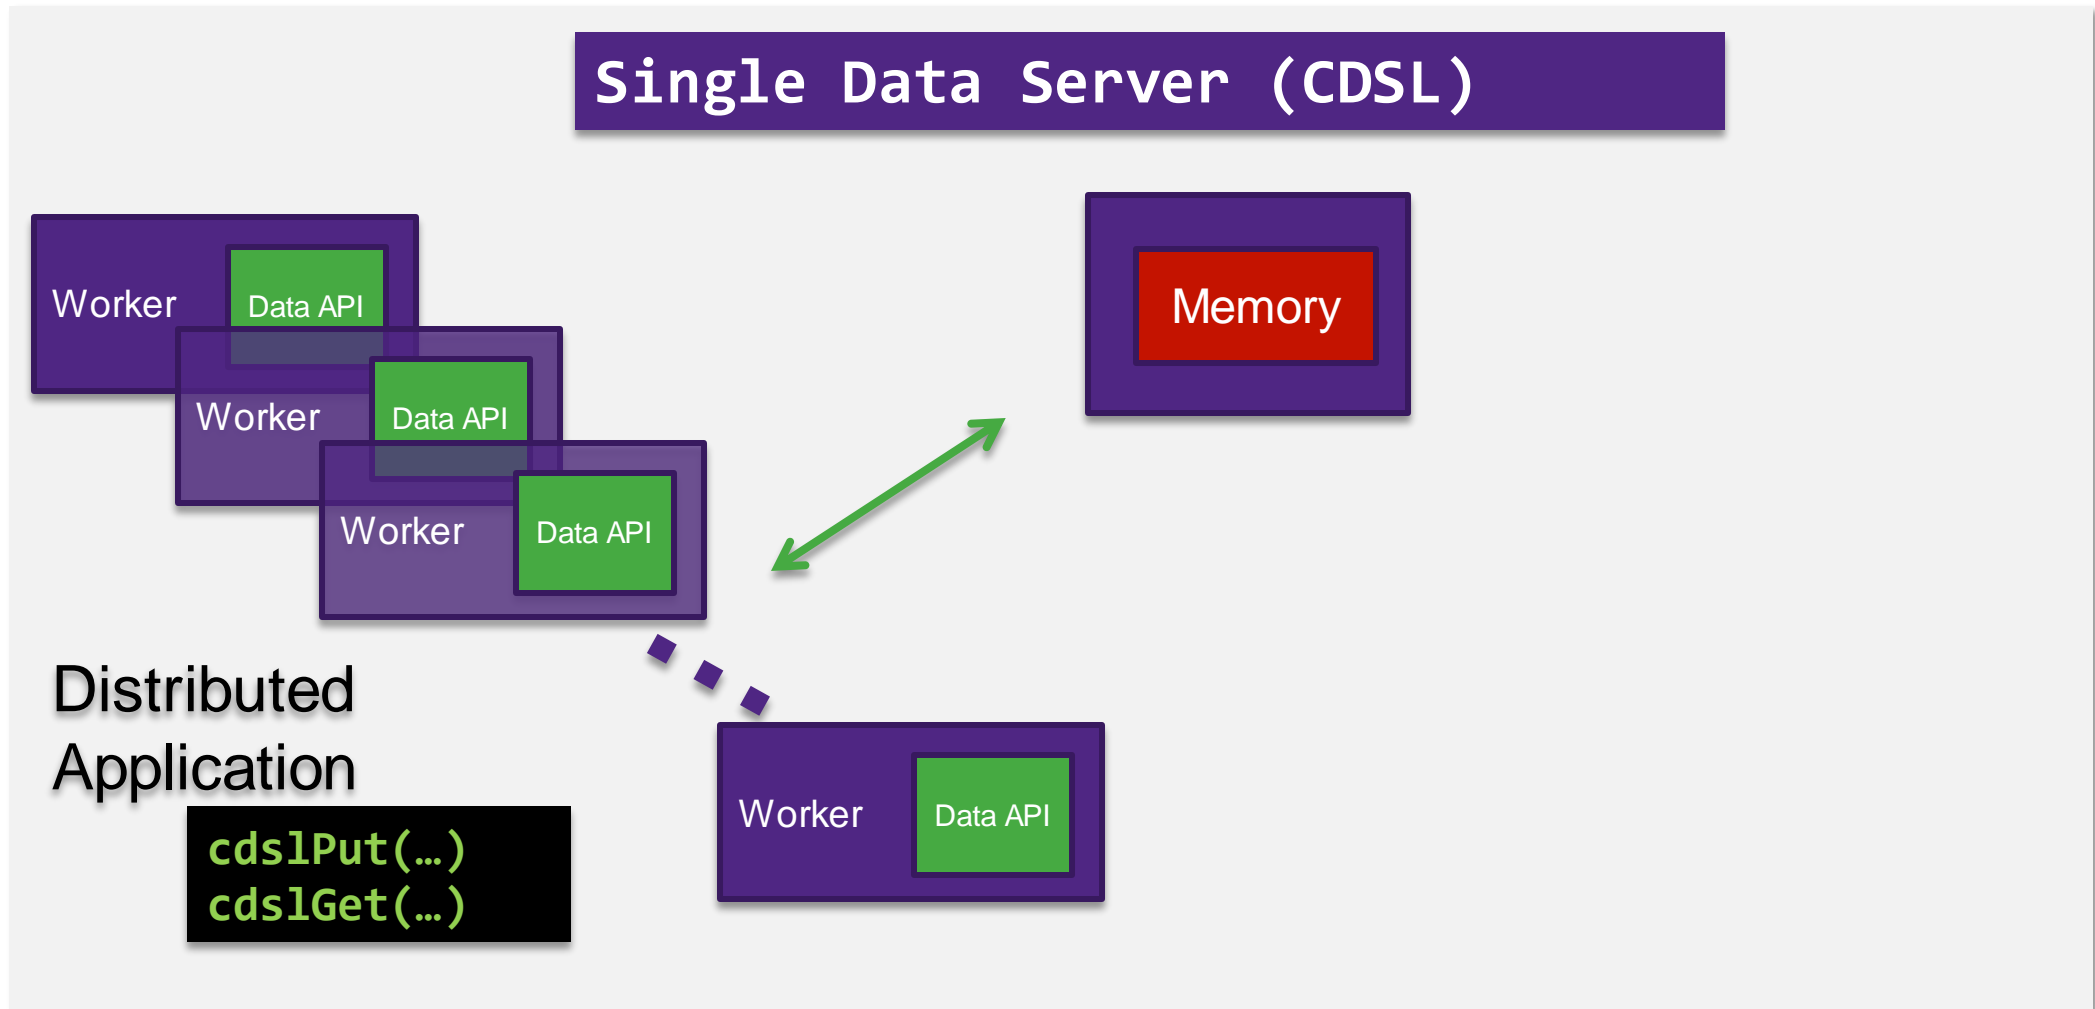


CDSL – providing a data management backplane



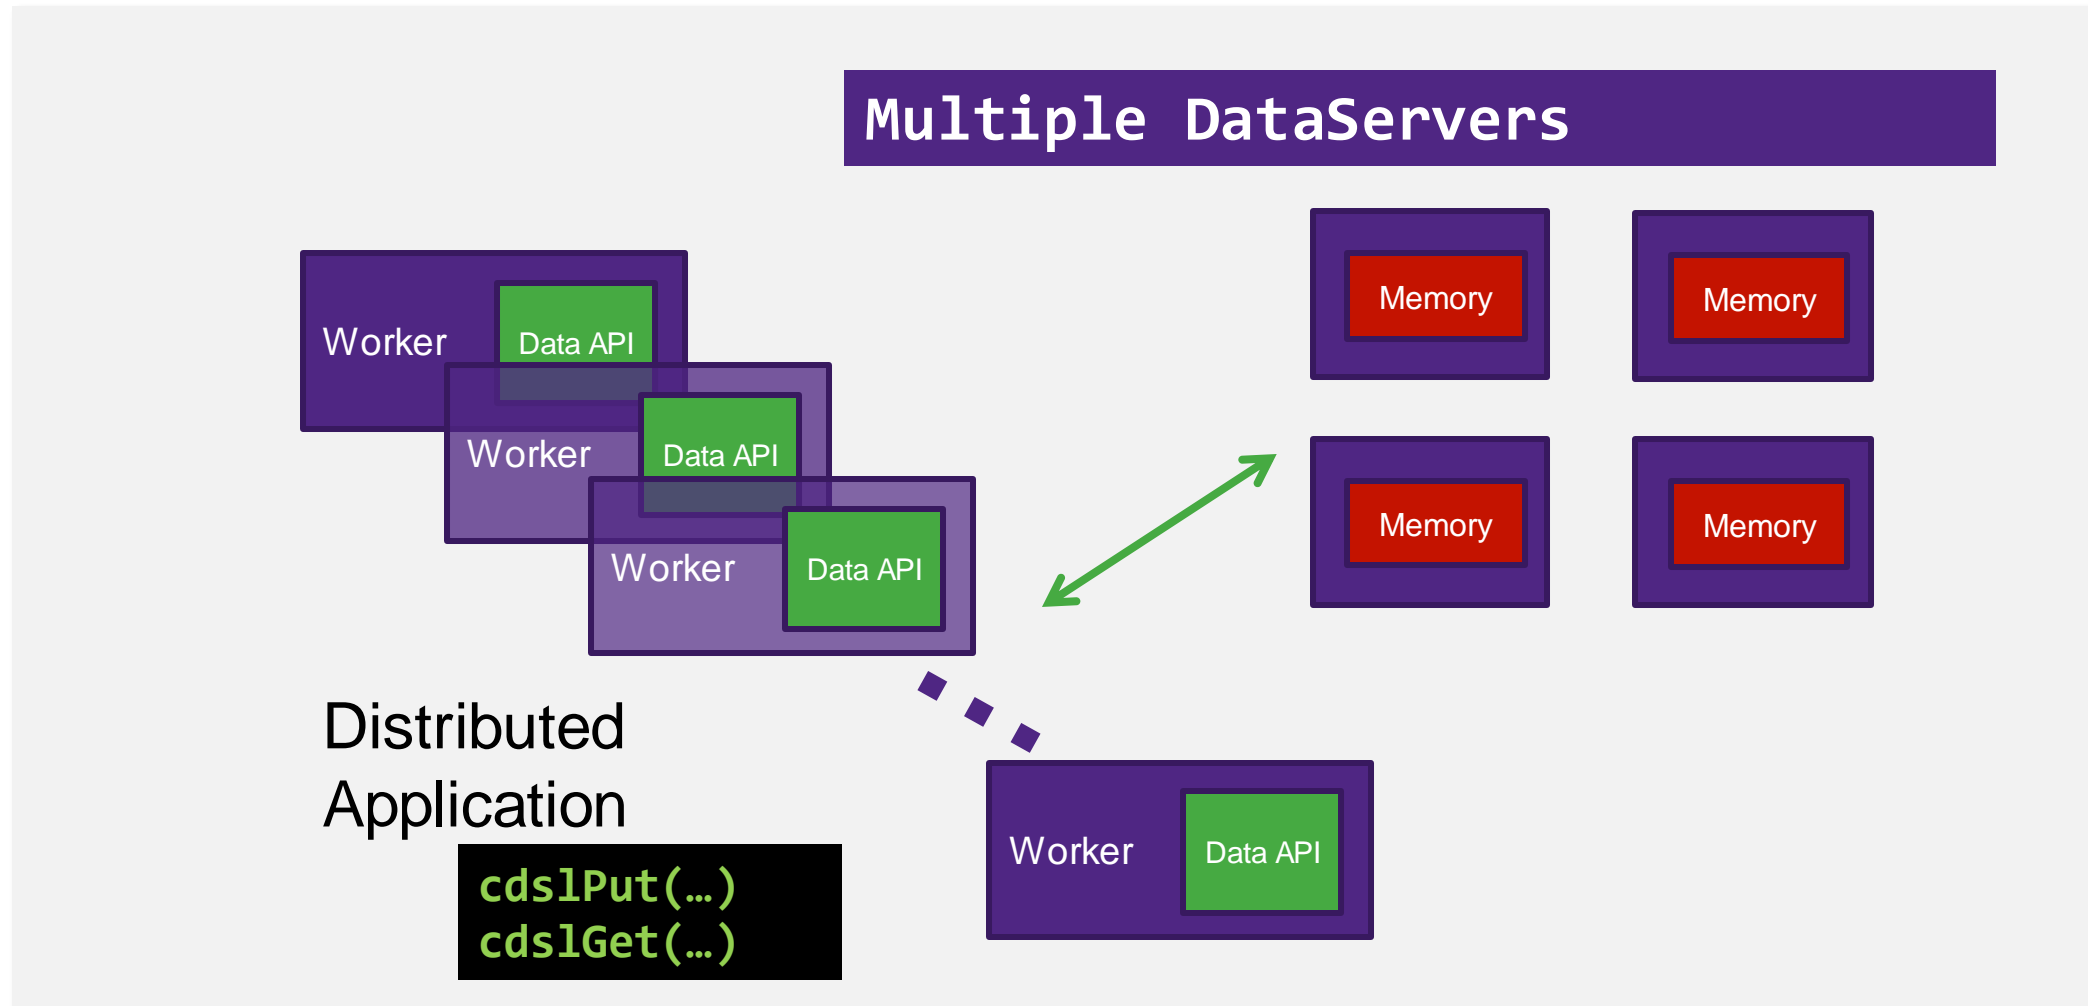
CDSL use model: single instance

Start and use one in-memory Key-Value Store



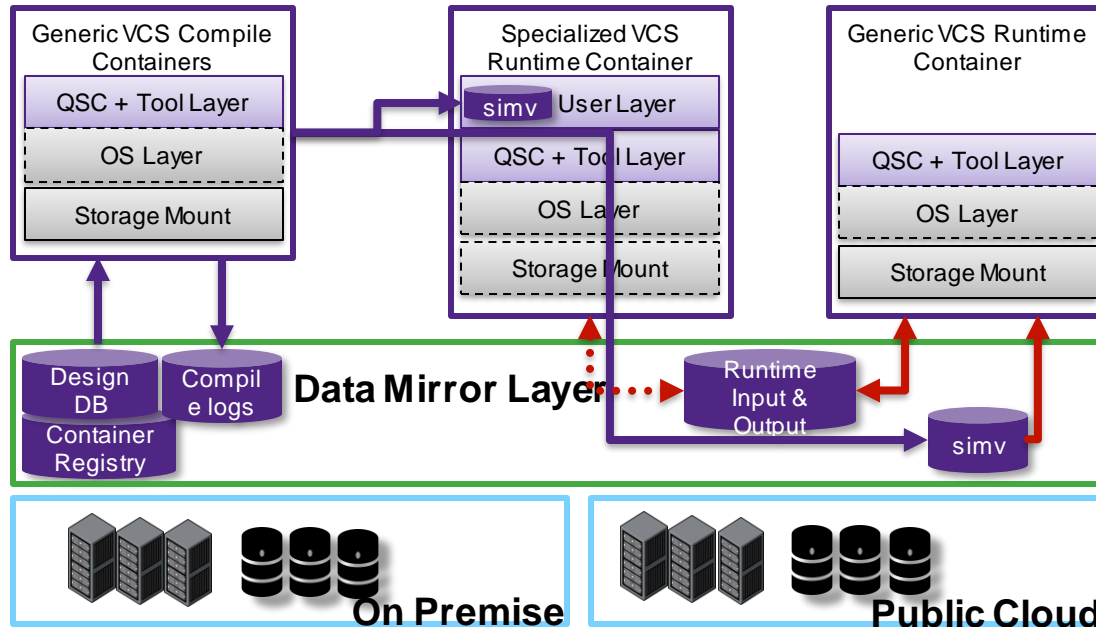
Horizontally scaled CDSL Data Servers

Shard data across a pool of in-memory Key-Value Stores



Case Study 1 - VCS

VCS – Data Consistency / Transfer



• Future Work

- Eliminate NFS bottlenecks on cloud
 - Read only partitions backed by object store with a Posix interface
- Sync of summarized ML data on cloud and on-prem
- Distributed constraints caching

VCS Container Use Model

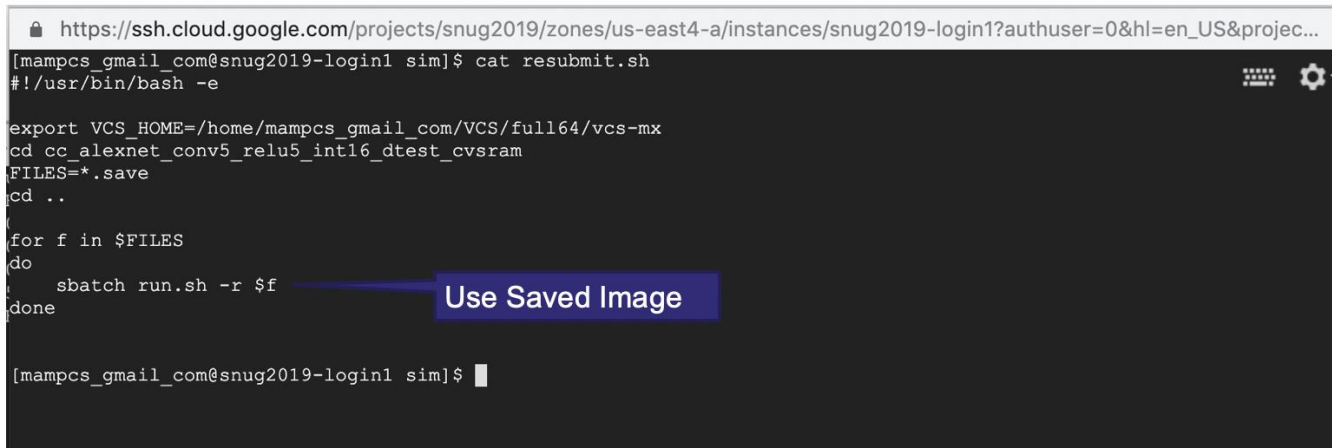

- `vcs --container=singularity[:path to base image] --container_include=<file name>`
 - If path to base image is not specified, default image is used from the vcs installation
 - File name is the path to a file that contains a list of directories/files to be copied into the container image
- `vcs --container=docker:[from_registry/]image_name[:version] --container_include=<file name>`
`--container_publish=to_registry`
 - `from_registry` is the details of the registry where the base container image is available
 - `to_registry` is the details of the registry when the built image should be published
 - File name is the path to a file that contains a list of directories/files to be copied into the container image

VCS - Spot Instances

Shutdown Script

VCS Auto Save

Relaunch Saved Simulation



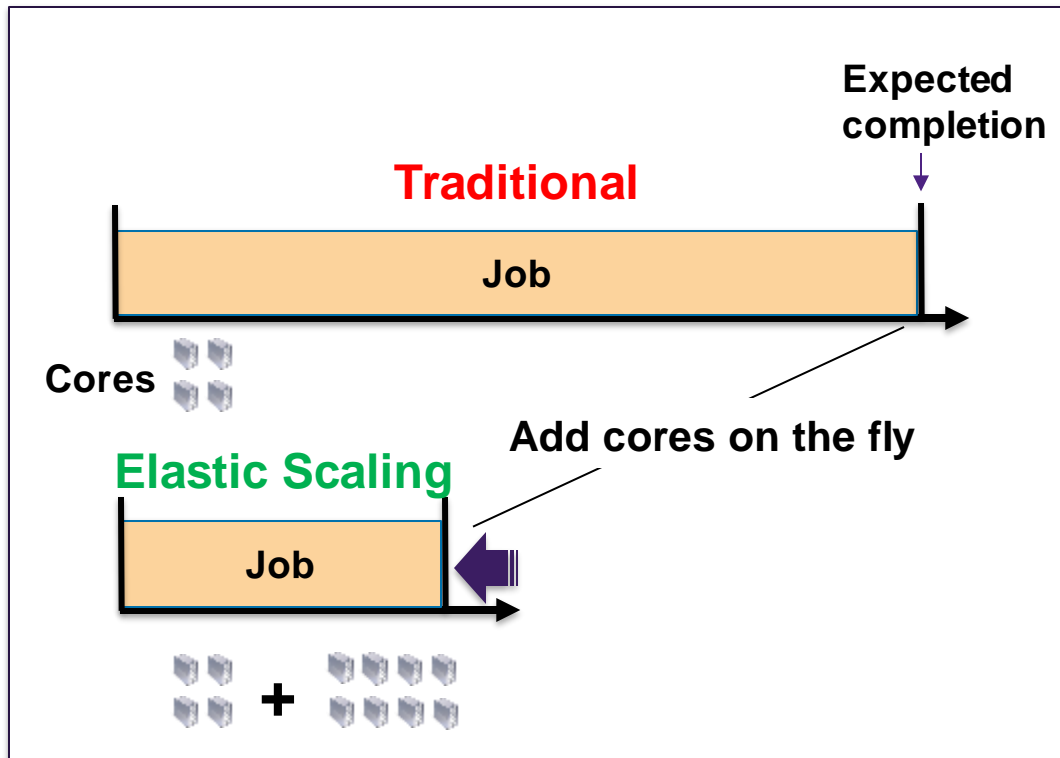
- Future Work

- Support for other cloud providers
- Tighter integration with schedulers for restart
 - LSF
 - Univa
- Periodic saves
- Support for heterogeneous instance types
- Selection of spot instances based on spot pricing and simulation job requirements

Case Study 2 – ICV

ICV - Elastic Scalability Technology for the Cloud

Add cores on the fly to a running job



Example: IC Validator Elastic Scalability

